

# AES-OTR v2

Kazuhiko Minematsu (NEC Corporation)  
Joint work with Maki Shigeri and Hiroyasu Kubo  
(NEC Solution Innovator)

# AES-OTR (v2)

- OTR blockcipher mode of operation (Eurocrypt 2014), defined with AES
- Nonce-based AE (NAE)

## Features:

- Rate-1 (needs one AES call for one block)
  - Parallelizable for encryption/decryption
  - Inverse-free
- 
- Unique AES-based NAE candidate achieving all of them

# Limitations

- NAE: nonce must be unique
  - Security guarantee is the same as current standards (CCM/GCM)
  - Assumption : AES = PRP (or PRF)
- No protection against nonce-misuse and decryption-misuse
- Theoretical limit speed and size : AES in counter mode

# A change from v1

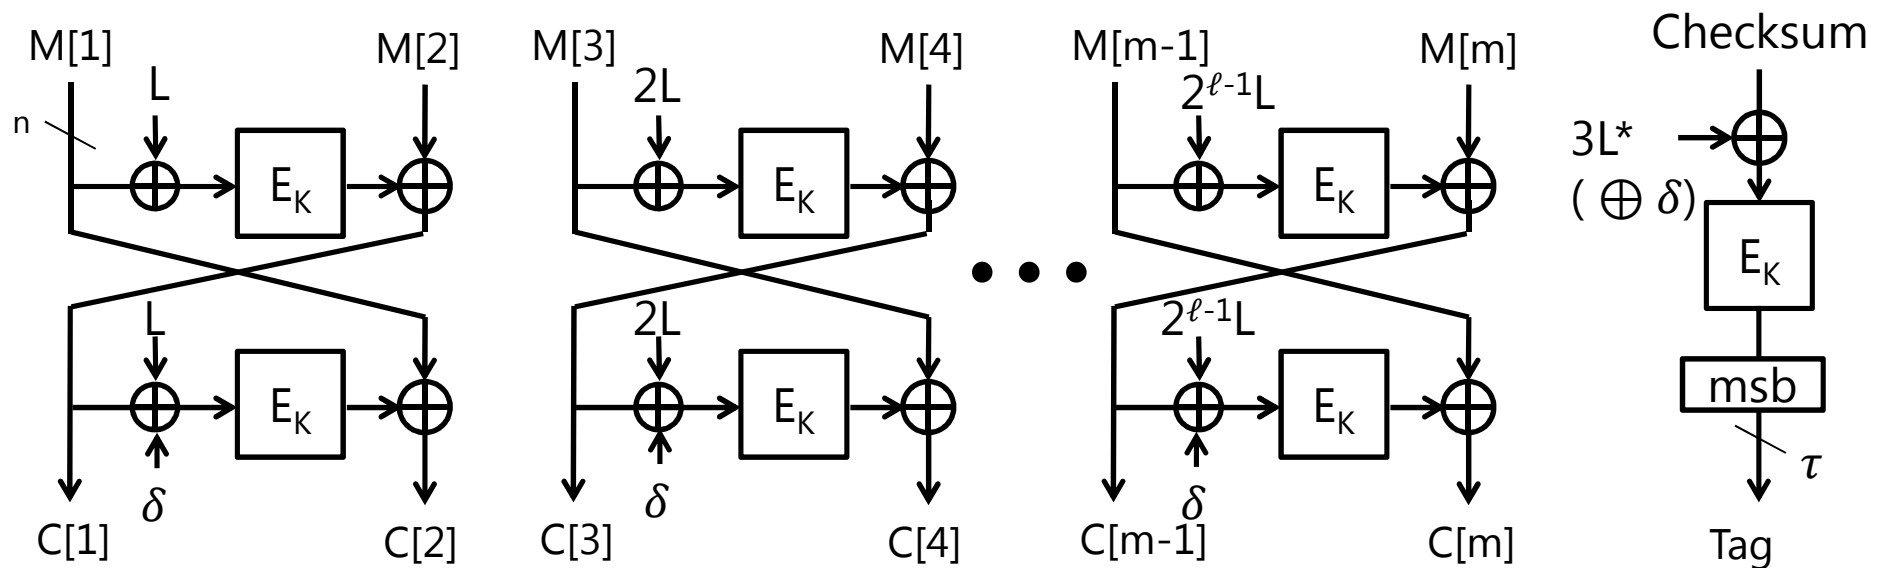
- A small change in nonce processing:
- Tag length ( $\tau$ ) is encoded with nonce encryption
- $\text{Format}(N, \tau) = (\text{len}(\tau) \bmod n)_7 || 00..01 || N$ 
  - The same function as OCB
  - Not allowing variable-tag-length ( $\tau$  is a parameter fixed in use)
  - For preventing trivial misuse (e.g. chopping tag more than what agreed)
  - No performance penalty

# Processing of OTR

- Use two-round Feistel permutation with XE-mode as round function
  - Using  $GF(2^{128})$  doubling
- Two-round Feistel can work as double-block cipher in OCB-like AE
  - Checksum is the sum of even blocks
- Last one or two blocks are processed a variant of two-round or CTR to avoid expansion
- Two AD processing versions (PMAC/CMAC) with different points to add MAC output

# Basic structure of OTR (no AD)

- $\delta = E(\text{Format}(N, \tau))$
- $L = 4 \delta$ 
  - See v2 doc for details



# Software performance analysis : the case of AESNI


# AES-OTR in AESNI

- AESNI: AES hardware instruction available in modern high-end CPUs (Intel, AMD)
- Basic expectation: AES-OTR runs fast, but slightly slower than AES-OCB
  - According to Bogdanov-Lauridsen-Tischhauser (ePrint 2014, FSE 2015) : the difference in peak speed (i.e. for long messages) : 0.15~0.2 cycles/byte
- A natural tread-off
  - inverse-freeness can be useful for many other platforms (e.g. hardware, embedded CPUs)
  - but closing the gap is good anyway



# AES-OTR in AESNI

- One possible reason: GF maskings
- 1<sup>st</sup> Feistel round: GF doubling
- 2<sup>nd</sup> Feistel round: additional XOR of  $\delta$ 
  - Which can be absorbed into the first round key, by precomputing  $RK_1 \oplus \delta$
- OCB (of CAESAR candidate) does not need doubling in motion ; gray code with precomputed values
  - At the cost with increased memory



# GF doubling

- $X \rightarrow (X \ll 1) \oplus \text{msb}(X) \cdot (0x87)$ 
  - we also need endian conversion, which can be done in one instruction (pshfub)
- We usually perform doubling over 128-bit XMM register
- But there is no single instruction to perform one-bit shift of XMM
- A common method needs 5 instructions
  - OTR uses a doubling for each 32 bytes  $\rightarrow$  overhead is around 0.15 c/b
    - Can be a major factor
- We review known doubling methods

# Krovetz

- (in optimized OCB C code\*1)
- 5 instructions
  - A clever way to compute carry by arithmetic right shift (`mm_srai_epi32`)
  - All fast instructions

Example:

```
void doubling(__m128i in, __m128i *out) {  
    const __m128i mask = _mm_set_epi32(135, 1, 1, 1);  
    __m128i tmp = _mm_srai_epi32(in, 31);  
    tmp = _mm_and_si128(tmp, mask);  
    tmp = _mm_shuffle_epi32(tmp, _MM_SHUFFLE(2, 1, 0, 3));  
    *out = _mm_slli_epi32(in, 1);  
    *out = _mm_xor_si128(*out, tmp);  
}
```


\*1 available from (<http://www.cs.ucdavis.edu/~rogaway/ocb/>)

# Aoki-Iwata-Yasuda (DIAC 2012)

- Use movemask (MOVMSKPD) instruction for extracting two MSBs in two 64 bits
- 3 instructions + 1 table lookup
- Table lookup can be costly, MOVMSKPD is not the fastest (latency 3 @ Haswell)


Example:

```
void doubling_aoki(__m128i in, __m128i *out){  
    unsigned bm;  
    __m128i tbl[4];  
    tbl[0]= _mm_set_epi32(0x00,0x00,0x00,0x00);  
    tbl[1]= _mm_set_epi32(0x00,0x01,0x00,0x00);  
    tbl[2]= _mm_set_epi32(0x00,0x00,0x00,0x87);  
    tbl[3]= _mm_set_epi32(0x00,0x01,0x00,0x87);  
    bm = _mm_movemask_pd((__m128d)in);  
    *out = _mm_add_epi64(in, in);  
    *out = _mm_xor_si128(*out, tbl[bm]);  
}
```



# Aoki (SCIS 2013 (in Japanese))

- Batch doubling :  $X \rightarrow (2X, 4X, 8X, \dots)$ 
  - Useful for many schemes, incl. OTR
- For  $X = (Xa, Xb)$  (each 64 bits),
- 8 16-entry 16-byte tables storing  $Xa$ 's carry (shift-and-xor of 0x87) and  $Xb$ 's carry
  - Table lookups for  $Xa[0-3]$ ,  $Xa[4-7]$ ,  $Xb[0-3]$ ,  $Xb[4-7]$
  - Total 2K bytes
- After address calculation, each doubling can be done with **3** fast instructions with **2** table lookups



# Other methods

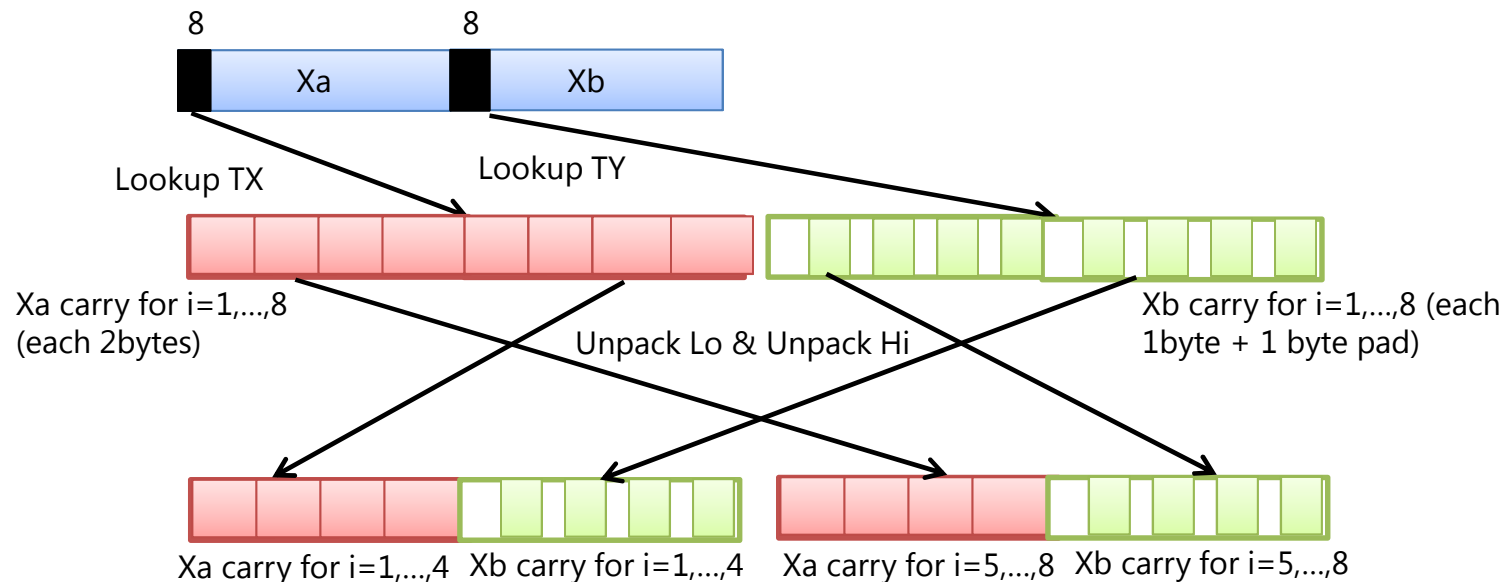
- Bogdanov et al. (eprint 2014, FSE 2015)
  - Similar to Krovetz
  - Investigated various options for implementing "if(MSB(X)==1)"
  - 4 instructions + "if(MSB(X)==1)"
- Chakraborty-Sarkar (ePrint 2014)
  - Survey of known methods and comparison of doublings of different fields
- Our first try
  - (64-bit x 2) version of Krovetz
  - Replace `_mm_srli_epi32` + `_mm_and_si128` with `_mm_maskload_epi64` (VPMASKMOVD in AVX2)
  - 4 instructions, but maskload is slow (latency 4, reciprocal t'put 2 @ Haswell), according to Agner Fog

# A quest for fast doubling

- We implemented Aoki, Batch Aoki, and Bogdanov et al.
- In our experiment, none of them constantly beat the original Krovetz
  - Except Aoki for 8 batch doublings, which runs almost as fast as Krovetz, and for some cases, it is slightly faster
  - Gain depends on microarchitecture and compiler

# A new batch method for 8 doublings

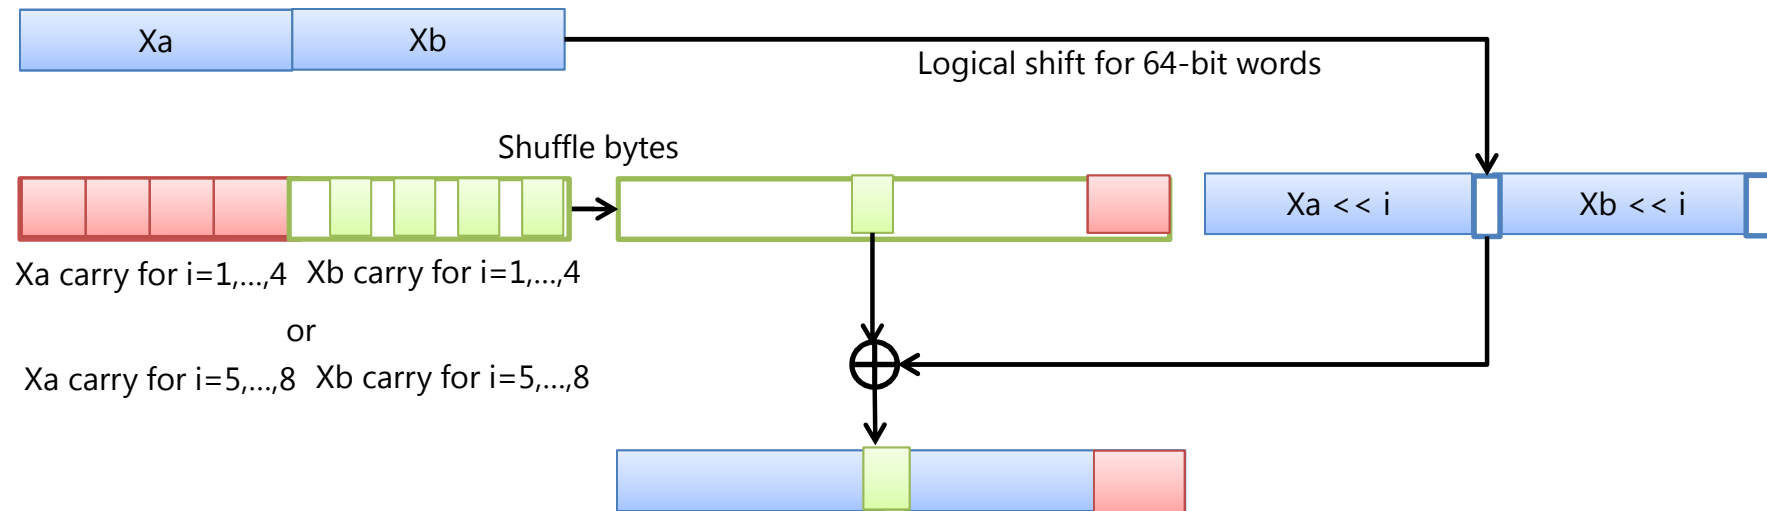
- Take xmm X as two 64 bits, (Xa, Xb)
- Use two 256-entry 16-byte tables (TX, TY) storing the results of Xa carry (2 bytes) and Xb carry (1 byte + 1 byte pad) for  $i=1,\dots,8$ -times doublings
  - total 8K bytes
- Swap upper and lower halves of TX and TY 2 using 2 unpack instructions
- Two table lookups and two unpacks





# A new batch method

- Then for each  $i$ , a doubling ( $2^i X$ ) can be done by a shuffle instruction (to obtain the  $i$ -th offset), a shift instruction, and a XOR instruction
- Thus **3** fast instructions



# Effect

- We observed improvement from Krovetz
- For example on Haswell Core i7 with gcc 5.1.0
  - (a batch version of) Krovetz and batch Aoki : 4.8
  - The new method : 4.1
  - If with endian conversions, the gain is around 0.5 ~ 1.3 c/b, depending on microarchitecture and compiler
- The same routine can be useful for other CAESAR candidates, such as
  - COPA, ELmD, POET (in MAC part), SHELL
  - Non-CAESAR : PMAC, XTS etc.
- We need care when side-channel analysis against this table lookup is a practical risk

# AES-OTR in AES-NI

- The code is written with intrinsic and new batch doubling
- ~10% improvement from [BLT15] @ 2Kbyte message

Haswell Core i7-4770 CPU @ 3.40GHz (CentOS gcc 5.1.0),  
AES-OTR w/ AES-128, no AD

Msglen (byte)	Enc (median, c/b)	Dec (median, c/b)
1024	1.02	1.11
2048	0.84	0.86
4096	0.78	0.79

# AES-OTR in ARMv7

# AES-OTR in ARMv7

- Use Bitslice AES available from SUPERCOP (originally Kasper-Schwabe ,CHES 2009)
  - Done 8 blocks in parallel
- Also single-block AES (standard T-table)
  - Nonce/tag encryption
- Use NEON SIMD engine, intrinsic
- No optimization wrt doubling
- Platform: Beaglebone Black (Cortex-A8 1GHz), with gcc 4.7.3

# Results

- Peak speed : ~23.5 c/b ( +7% of AESBS)
- For reference, in Gouvea-Lopez (CTRSA 2015) GCM runs 32.8 c/b, using BS-AES, on Cortex A9
- Wanted : fast constant-time single-block AES on ARMv7 (vector-permutation-based one) for gaining performance for short messages

AES-OTR (AES-128, no AD)

Msglen (byte)	Enc (median c/b)	AESBS ratio	Dec (median c/b)	AESBS ratio
1056	25.42	1.14	25.42	1.14
2080	24.19	1.07	24.2	1.07
16416	23.5	1.07	23.51	1.07

AES Bit-slice

Msglen (byte)	Enc (med c/b)
1056	22.24
2080	22.58
16416	21.87

AES T-table

Msglen (byte)	Enc (med c/b)
1056	41.31
2080	43.56
16416	43.54

# AES-OTR in FPGA (preliminary result)

# AES-OTR in FPGA

- A basic comparison of OTR/OCB/GCM
  - Reference versions (full-spec, unoptimized)
  - All modes use the same AES core (round-based)
  - Implement OTR and OCB. GCM mode taken from OpenCore (no decryption routine...)
  - Additionally, (experimental) dual-core implementation for OTR
- Environments: Altera Cyclone IV (EP4CE30F29C6), Quartus 13.1.4
- Many thanks to Tomonori Iida for implementation



# Results

## Size

		Logic Cell (AES)	LUT	Register
OTR parallel	ENC	6,647 (3,138)	6,618	1,576
	DEC	7,069 (3,120)	6,927	1,845
	ENC/DEC	8,056 (3,141)	7,901	1,846
OTR parallel AES2 Core	ENC	12,682 (5,483)	12,636	2,392
	DEC	13,041 (5,587)	13,007	2,391
	ENC/DEC	15,773 (5,633)	15,736	2,394
OTR serial	ENC	6,322 (3,144)	6,282	1,448
	DEC	6,814 (3,103)	6,693	1,717
	ENC/DEC	7,631 (3,122)	7,496	1,718

		Logic Cell	LUT	Register
OCB	ENC	7,447 (3,141)	7,357	1,522
	DEC	10,657 (3,053/3,100)	10,657	1,917
	ENC/DEC	11,711 (3,123/3,144)	11,515	1,923

		Logic Cell	LUT	Register
GCM	ENC	7,167 (3,127)	6,973	1,434

## Speed

For 12-byte  
Nonce, 16-  
byte AD, 32-  
byte Plaintext

		Clock cycles	Max Freq. (MHz)	Time (ns)
OTR parallel	ENC	72	106.58	676
	DEC	72	108.95	661
	ENC/DEC	72	108.31	665
OTR parallel AES 2 Core	ENC	39	103.63	376
	DEC	49	108.98	450
	ENC/DEC	39/49	96.59	----
OTR serial	ENC	72	105.30	684
	DEC	72	105.83	680
	ENC/DEC	72	104.99	----

		Clock cycles	Max Freq. (MHz)	Time (ns)
OCB	ENC	72	110.98	649
	DEC	72	99.45	724
	ENC/DEC	72	97.90	735

		Clock cycles	Max Freq. (MHz)	Time (ns)
GCM	ENC	80	104.18	768

# Rough summary

- OTR is the smallest (in particular, serial ADP is small)
- OTR & OCB run with the same cycles, have similar speed
- If AD is always empty, -1,000 LE is possible

		Logic Cell	LUT	Register
OTR serial, AD empty	ENC	5,640(3,121)	5,601	1,444

- OTR's gain from GCM is moderate
  - As expected: OTR focus the balanced performance on Sw/Hw, from low-end to high-end
- Need more studies...

Thank you!